

DISTRIBUTED SYSTEM DESIGN

MPI

What is MPI?

- MPI = Message Passing Interface
- MPI is a library of routines that can be used to create parallel programs.

Fundamentals: *Communicators & Groups*

- MPI defines **communicators and groups** to define which collection of processes may communicate with each other
- Most MPI routines/functions require a *communicator* as an input parameter
- For simplicity, we'll be using the **MPI_COMM_WORLD** communicator
 - This communicator includes *all of your MPI processes*

Fundamentals: *Ranks*

- Within a communicator, each process has its own and **unique ID** or *rank*
 - These IDs are commonly used conditionally to control program execution
- Ranks start from **0**

MPI Routines

- `MPI_Init(int *argc, char ***argv)`
- This initializes the MPI execution environment.
 - Therefore, this **must** be called (once) **at the start of every MPI program**

MPI Routines

- `MPI_Comm_size(MPI_Comm comm, int *size)`
- This determines the number of processes in the group associated with the `comm` communicator

MPI Routines

- `MPI_Comm_rank(MPI_Comm comm, int *rank)`
- This determines the **rank** of the calling process within the **comm** communicator.

MPI Routines

- **MPI_Wtime()**
- This returns an elapsed wall clock time in seconds (double precision) on the calling processor.
 - We'll use this to **measure the runtime** of an MPI program

MPI Routines

- `MPI_Send(void *buf, int count,
MPI_Datatype datatype, int dest, int tag,
MPI_Comm comm)`
- This is a basic **blocking send** operation. It returns only after the application has sent the data to the recipient(s)

MPI Routines

- `MPI_Recv(void *buf, int count,
MPI_Datatype datatype, int src, int tag,
MPI_Comm comm, MPI_Status *status)`
- This receives a message and blocks until the requested data is available in the application buffer

MPI Routines

- **MPI_Finalize()**
- This terminates the MPI execution environment.
 - This should be called **at the end** of every MPI program

Running MPI

MPI Examples

- Together, we'll program two MPI examples:
 - HelloWorld
 - A Distributed Sum Program